

# Linear Models and ML Fundamentals II

ML4SE

Denis Litvinov

September 21, 2021

# Table of Contents

## 1 Binary Classification

- Losses
- Logistic Regression
- SVM
- Robustness
- Quality Metrics

## 2 Model Complexity

## 3 Multiclass Classification

- One vs Rest

- One vs One

- Multinomial

- Softmax

- Quality Metrics

## 4 Validation

- Hold Out

- Cross Validation

- Leave One Out

- Common Pipeline

# Recap on Choice of Quality Metric



# Problem Statement



Given dataset  $\{(x_i, y_i)\}_{i=1}^N$  of i.i.d. objects

Or, equivalently given:

$X \in R^{N \times D}$  - feature matrix, where  $D$  is dimension of feature space  
and  $N$  - number of objects.

$Y \in \{0, 1\}^N$  - target vector

Sometimes we will use notation  $Y \in \{-1, 1\}^N$



# Bayesian Classifier

Bayesian classifier is the best possible classifier given we know all joint distribution  $P(x, y)$  of features and labels. (Which is an unrealistic assumption.)

Bayesian risk:

$$R = \sum_{x,y} I[h(x) \neq y] P(x, y) c_y$$

, where  $c_y$  is cost function for misclassification, e.g.  $\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$  or  $\begin{pmatrix} 0 & 0.7 \\ 0.3 & 0 \end{pmatrix}$

function  $h(x)$  - decision function.


Which is minimized by

$$h(x) = \arg \max_y P(y|X)$$

$$h(x) = \arg \max_y P(X|y) P(y) c_y$$

# Bayesian Classifier

⇒ Cost function and prior class probabilities are interchangeable!




# Margin

For binary classification with  $y \in \{-1, 1\}$  a variable  $z = yh(x)$  is called **margin**.

Positive margin corresponds to successful classification, negative margin corresponds to error.

$|yh(x)|$  is a distance to decision boundary, which can be interpreted as confidence in classification of the object.



# Logistic Regression

Suppose  $\hat{y} = h(x)$  and  $y \in \{0, 1\}$ .

Show, that  $h(x_i)$  should be  $p(y = 1|x_i)$ .

Probability to generate such samples from the point view of  $h(x)$  is a likelihood  $L$ :

$$L = \prod_{i=1}^N h(x_i)^{[y_i=1]} (1 - h(x_i))^{[y_i=0]} \rightarrow \max_h$$

$$\log L = \sum_{i=1}^N [y_i = 1] \log h(x_i) + [y_i = 0] \log(1 - h(x_i)) \rightarrow \max_h$$

$$-\log L = - \sum_{i=1}^N [y_i = 1] \log h(x_i) + [y_i = 0] \log(1 - h(x_i)) \rightarrow \min_h$$

$$E[-\log L|x] = -p(y = 1|x) \log h(x) - p(y = 0|x) \log(1 - h(x)) \rightarrow \min_h$$



# Logistic Regression

$$\frac{\partial E[-\log L|x]}{\partial h} = -\frac{p(y=1|x)}{h(x)} + \frac{1-p(y=1|x)}{1-h(x)} = 0$$

$$\implies h(x) = p(y=1|x) = \sigma(w^T x)$$

What are requirements for  $\sigma(x)$  ?


# Logistic Regression

Requirements for  $\sigma(x)$ :

- 1  $\forall x \in \mathbb{R}^D \sigma(x) \in [0, 1]$
- 2  $\sigma(0) = 0.5$
- 3  $\sigma(-x) = 1 - \sigma(x)$
- 4  $\sigma$  - non-decreasing, continuous and differentiable

Sigmoid

$$\sigma(z) = \frac{1}{1+e^{-z}}$$



# Logistic Regression

You can find 2 different formulae for logistic loss: via cross-entropy as shown above

$$Loss(y_i, p_i) = -y_i \log p(y_i = 1|x_i) - (1 - y_i) \log(1 - p(y_i = 1|x_i))$$


where probability of  $y = 1$  class given sample  $x$  is

$$p(y = 1|x) = \sigma(w^T x) = \frac{1}{1 + e^{-w^T x}}$$


and with margins

$$Loss(y_i, x_i) = \log(1 + e^{-y_i w^T x_i})$$

# Logistic Regression



# Support Vector Machines



# Support Vector Machines

Suppose we have some linear decision surface  $h(x) = \text{sign}(w^T x)$   
Then distance from point  $x_0 \in R^D$  to decision surface is

$$\rho(x_0, h) = \frac{|w^T x|}{\|w\|_2}$$

Choose scale of  $w$  such that

$$\min_x |w^T x| = 1$$

Then distance from decision surface to the nearest object is

$$\min_x \frac{|w^T x|}{\|w\|_2} = \frac{1}{\|w\|_2} \min_x |w^T x| = \frac{1}{\|w\|_2}$$

For linear separable case we have optimization problem:

$$\begin{cases} \frac{1}{2} \|w\|_2^2 \rightarrow \min_w \\ y_i w^T x_i \geq 1 \end{cases}$$

# Support Vector Machines

For linear inseparable case we introduce corrections for each object

$\xi_i$ :

$$\begin{cases} \frac{1}{2} \|w\|_2^2 + C \sum_{i=1}^N \xi_i \rightarrow \min_{w, \xi_i} \\ y_i w^T x_i \geq 1 - \xi_i \\ \xi_i \geq 0 \end{cases}$$

OR:


$$\xi_i = \max(0, 1 - y_i w^T x_i)$$

Then,

$$Loss = \frac{1}{2C} \|w\|_2^2 + \sum_{i=1}^N \max(0, 1 - y_i w^T x_i) \rightarrow \min_w$$

Unlike logistic regression, weight norm penalty already build in the model.

# Losses







# Robustness

Model is called robust if its performance do not change significantly for new samples drawn from the same distribution  $P(x, y)$ .

In other way, model robustness depends on how the model handle the outliers.



# Adversarial Examples


Adversarial example is an object from original  $P(x, y)$  with small perturbation in  $x$  such that it makes the model to give false predictions.



**"panda"**

57.7% confidence

+  $\epsilon$



=



**"gibbon"**

99.3% confidence

# F1 score

$$\text{Accuracy } acc(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N [\hat{y}_i = y_i]$$

		Predicted class	
		<i>P</i>	<i>N</i>
Actual Class	<i>P</i>	True Positives (TP)	False Negatives (FN)
	<i>N</i>	False Positives (FP)	True Negatives (TN)

$$\text{precision} = \frac{TP}{TP + FP}$$

$$\text{recall} = \frac{TP}{TP + FN}$$

$$F1 = \frac{2 * \text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$


# AUC

$FPR = \frac{FP}{FP+TN}$  false positive rate

$TPR = \frac{TP}{TP+FN}$  true positive rate

$AUC$  = area under the curve  $ROC$

$ROC(t) = (TPR(t), FPR(t))$  is parametrized by threshold  $t$  on the probability  $p(y = 1|x)$



How does class imbalance affect quality metrics given above?

# Empirical Loss Minimization

In general, we want to optimize Expected Risk:

$$R = E[Loss(x, y)] = \int_{-\infty}^{\infty} Loss(x, y) dP(x, y) = Pr_{(x_i, y_i) \sim D}[Loss(x_i, y_i)]$$

But since we don't know the joint distribution  $P(x, y)$ , we can only deal with Empirical Risk (Loss functional):

$$\hat{R} = \sum_{i=1}^N Loss(x_i, y_i)$$

# Generalization error

How well does  $\hat{R}$  approximates  $R$ ?

Using Hoeffding's inequality it can be shown, that given  $N$  random examples, and  $\forall \delta > 0$ , with probability  $Pr \geq 1 - \delta$ , the following upper bound holds on the generalization error of  $h$ :

$$R \leq \hat{R} + \sqrt{\frac{\ln(1/\delta)}{2N}}$$

For finite hypothesis space  $H$  under the same conditions:

$$R \leq \hat{R} + \sqrt{\frac{\ln |H| + \ln(1/\delta)}{2N}}$$

where  $\ln |H|$  serves as a measure of model complexity.

# Rademacher Complexity

How to approximately estimate model complexity?

Given a dataset  $\{(x_i)\}_{i=1}^N$ ,

Let  $\sigma_i \sim \text{random}(\{-1, 1\})$  be random labels for each  $x_i$ .

Consider binary classification task. We expect a model to have high complexity if it can fit well any random label assignment on the dataset.

Rademacher Complexity can be estimated by expected average margin


$$E_{\sigma} \left[ \max_{h \in H} \frac{1}{N} \sum_{i=1}^N \sigma_i h(x_i) \right]$$

We also say, that model with high complexity tends to memorize training samples.

# Data Augmentation

Generalization error depends on number of samples  $N$  in the dataset.  
How to increase  $N$ ?

Apply some transformation to original data which is invariant to some desirable properties.





# One vs Rest


Idea: build multiclass classifier from several binary classifiers

Train  $K$  binary classifiers.

$$\hat{y} = \arg \max_k h_k(x)$$

Note:

- 1  $h_k$  is unbalanced even if initial problem was balanced
- 2 scale of the confidence values may differ between the binary classifiers  $b_k$




# One vs One

Idea: build multiclass classifier from several binary classifiers

Train  $K(K-1)/2$  binary classifiers.

$$\hat{y} = \arg \max_k \sum_{i \neq k} h_{ik}(x)$$

Note: One vs one is less prone to imbalance in dataset



# Multinomial

Cross-entropy loss:

$$Loss(y_i, p_i) = - \sum_{k=1}^K y_{ik} \log p(y_{ik} = 1 | x_{ik})$$

where  $k$  is a number of classes

$$y_{ik} = I[y_i = k]$$

Probability if  $x_i$  has class  $k$  (in vector form):

$$p(y_i = k | x_i) = \text{softmax}(x_i^T W)_k$$

where  $W \in R^{D \times K}$



# Softmax

Requirements:

- 1 non-negative
- 2 sums to 1 (is a probability)
- 3 monotonic increasing

for  $z \in \mathbb{R}^K$

$$\text{softmax}(z)_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$



# Quality Metrics

There is no direct quality metric, it is assembled from metrics for binary classification problems.

$$micro\_precision = \frac{\sum_k TP_k}{\sum_k TP_k + \sum_k FP_k}$$

$$micro\_recall = \frac{\sum_k TP_k}{\sum_k TP_k + \sum_k FN_k}$$

$$micro\_f1 = \frac{2 * micro\_precision * micro\_recall}{micro\_precision + micro\_recall}$$

weight in proportion of class size:

$$weighted\_precision = \frac{1}{K} \sum_k \frac{|K|}{N} \frac{TP_k}{TP_k + FP_k}$$

$$weighted\_recall = \frac{1}{K} \sum_k \frac{|K|}{N} \frac{TP_k}{TP_k + FN_k}$$

$$weighted\_f1 = \frac{2 * weighted\_precision * weighted\_recall}{weighted\_precision + weighted\_recall}$$

# Quality Metrics

Macro averaging is insensitive to class imbalance

$$macro\_precision = \frac{1}{K} \sum_k \frac{TP_k}{TP_k + FP_k}$$

$$macro\_recall = \frac{1}{K} \sum_k \frac{TP_k}{TP_k + FN_k}$$

$$macro\_f1 = \frac{2 * macro\_precision * macro\_recall}{macro\_precision + macro\_recall}$$

# Validation

Till now:

- 1  $\{x_i, y_i\}_{i=1}^N$  is sampled from  $P(x, y)$
- 2 Split into to non-overlapping subsets (train and test)
- 3 Model  $h(x_i; w, \theta)$  is described by its trainable weights  $w$  and non-trainable hyperparams  $\theta$
- 4 choose some hyperparam value  $\theta = \theta_0$  and train model

$$\sum_{i \in \text{train}} \text{Loss}(h(x_i; w, \theta_0), y_i) \rightarrow \min_w$$

- 5 test model performance on test dataset

$$\hat{R}(\theta_0) = \sum_{i \in \text{test}} \text{Loss}(h(x_i; w, \theta_0), y_i)$$

- 6 we expect that it is a good approximation of true generalization error

$$R(\theta_0) = E_{(x,y) \sim P(x,y)} [\text{Loss}(h(x_i; w, \theta_0), y_i)]$$

- 7 How to choose hyperparam  $\theta$ ?

# Validation

Usually want to optimize hyperparams by testing several values of  $\theta$  on the test set and choosing the best one.

But the performance on the test set  $\hat{R}(\theta)$  (empirical risk) is a random variable, which can depend on the particular train test split! Here validation comes into play.

We can say that  $\hat{R}(\theta)$  is a point estimate of expected risk  $R(\theta)$ , which has its bias and variance.

Different validation schemes try to minimize bias or variance or both.




# Hold Out

Given dataset of  $m$  objects, create  $m$  experiments:

- 1 create split train:val, usually in proportion 70:30, 80:20 or 90:10
- 2 fit model weights on train subset
- 3 evaluate performance on the val subset

Properties:

- High bias and low variance of estimate
- $O(1)$  complexity
- Usually done when we have large dataset and or very heavy model



# Cross Validation


$k$  = number of folds folds = non-intersecting subsets of the dataset

Make  $k$  experiments:

- 1 create split for  $k - 1:1$
- 2 train on  $(k - 1)$  folds and evaluate performance on the  $k$ -th fold
- 3 change split
- 4 Average scores over all experiments

Properties:

- Moderate bias and variance of estimate
- $O(k)$  complexity




# Leave One Out



Cross validation with  $k = N$

Properties:

- Low bias and High Variance of estimate
- $O(N)$  complexity
- Usually done when we have very small dataset
- There are performance metrics (e.g. AUC) that cannot be computed just on one sample.



# Hyperparam Search



# Common Pipeline

- 1 Split dataset for train, test parts
- 2 Choose validation scheme on training data
- 3 Train model on the train dataset without regularization, try to achieve zero training loss
- 4 Add regularization, tune hyperparams on validation
- 5 Evaluate final model performance on test dataset. Choose between different model families.

In practice we usually use chosen quality metric instead of loss function for choosing hyperparams and final testing.

# Common Pipeline

