

# Feed Forward Neural Networks

ML4SE

Denis Litvinov

October 6, 2021

# Table of Contents

**1** General Architecture

**2** Activation Functions

**3** Weight initialization

**4** Optimization

- Vanilla SGD

- SGD with momentum

- RmsProp

- Adam

- Learning Schedule

**5** Regularization


- Dropout

- BatchNorm

# General Architecture

NN as a composition of functions

$$F(x) = f_{w_n} \circ f_{w_{n-1}} \circ \dots \circ f_{w_1}(x)$$



# Composition of linear functions

$$F(x) = XW_1W_2$$

where  $X \in R^{N \times D_1}$  - features

$W_1 \in R^{D_1 \times D_2}$ ,  $W_2 \in R^{D_2 \times K}$  - weight matrices

$K$  - number of classes

# Neuron

$$y = \sum_{i=1}^N f(w_i x_i + b)$$


where  $f$  - some non-linear activation function

$w_i$  - learnable weights

$b$  - learnable bias, usually incorporated into  $X$

$y$  - output of neuron

$x$  - input of neuron



# Dense Layers


It's more convenient to express the same thing in vector form

$$Y = f(XW)$$

where  $X \in R^{N \times D_1}$  - input of layer

$Y \in R^{N \times D_2}$  - output of layer

$W \in R^{D_1 \times D_2}$  - learnable weight matrix



# Activation Functions


Why we do not use linear activations?

Activation functions supposed to be nice in the sense of gradient properties.

# Sigmoid

$$\sigma(z) = \frac{1}{1 + \exp^{-z}}$$


- vanishing gradient
- bad output distribution





# Tanh


- vanishing gradient



# RELU


$$RELU(z) = \max(0, z)$$

- dead neurons if  $z < 0$



# PRELU


$$PRELU(z) = \begin{cases} z, & \text{if } z \geq 0 \\ \alpha z, & \text{if } z < 0 \end{cases} \quad (1)$$



# ELU

$$ELU(z) = \begin{cases} z, & \text{if } z \geq 0 \\ \alpha(\exp^z - 1), & \text{if } z < 0 \end{cases} \quad (2)$$

- little longer computation than RELU



# Weight initialization

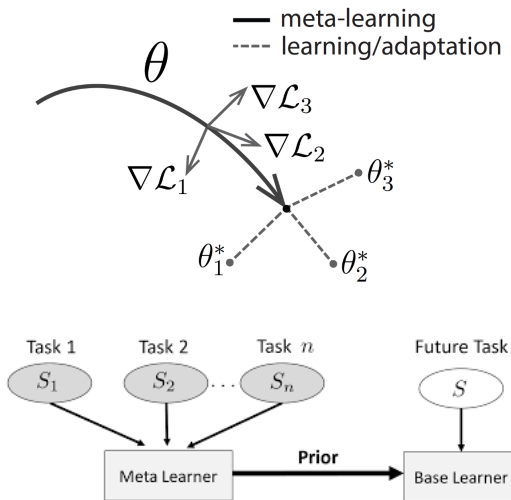
As we train out neural network with gradient descent, it is important to have good initial point to start. Usually use use:

- 1 Uniform distribution in  $[-d, d]$
  - 2 Normal distribution  $N(0, \sigma^2)$
  - 3 Xavier distribution (discuss later, in CNN)
- Why we use distributions centered around zero?
  - How it is connected with activation functions?
  - shared weights

# Transfer Learning


- 1 a big model is trained on a large dataset
- 2 learned weights from the model are used as a initialization for another, usually smaller, downstream task

# Meta Learning



# Loss surface


Many local minima





# Loss surface

Wide local minimum



# Vanilla SGD

$\theta_0 \leftarrow \text{init}$

for **random** batch on step  $t = 1..\text{max\_iter}$ :

$$\theta_t = \theta_{t-1} - \alpha \nabla_{\theta} J(\theta_{t-1})$$

$J$  - loss function

$\theta_t$  - learnable parameters at step  $t$

$\alpha$  - learning rate

- good theoretic properties
- slow convergence

# SGD with momentum

$\theta_0 \leftarrow \text{init}$

$m_0 \leftarrow 0$

for **random** batch on step  $t = 1..\text{max\_iter}$ :

$$m_t = \beta m_{t-1} + (1 - \beta) \nabla_{\theta} J(\theta_{t-1})$$

$$\theta_t = \theta_{t-1} - \alpha m_t$$


where  $m_t$  - accumulated gradient at step  $t$

$\beta$  - momentum parameter


# SGD with momentum

- Momentum cancels moves in "random" directions from stochastic nature of SGD
- Momentum inertia

SGD without momentum



SGD with momentum



# RmsProp

$\theta_0 \leftarrow \text{init}$   $v_0 \leftarrow 0$

for random batch on step  $t = 1..\text{max\_iter}$ :

$$g_t = \nabla_{\theta} J(\theta_{t-1})$$

$$v_t = \beta v_{t-1} + (1 - \beta) g_t^2$$

$$\theta_t = \theta_{t-1} - \frac{\alpha}{\sqrt{v_t} + \epsilon} g_t$$

where  $v_t$  - accumulated squared components of gradient

$\beta$  - parameter

$\epsilon \ll 1$  - to prevent division by zero

- gradient direction carries more information than its norm
- adjust gradient step size

# Adam

$\theta_0 \leftarrow \text{init}$   $v_0 \leftarrow 0$

$m_0 \leftarrow 0$

for random batch on step  $t = 1..\text{max\_iter}$ :

$$g_t = \nabla_{\theta} J(\theta_{t-1})$$

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$\theta_t = \theta_{t-1} - \frac{\alpha}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

where  $m_t$  - accumulated momentum


$v_t$  - accumulated squared components of gradient

$\beta_1, \beta_2$  - parameters

$\epsilon \ll 1$  - to prevent division by zero

# Adam


- essentially SGD with momentum + RmsProp
- corrections for  $\hat{m}_t$ ,  $\hat{v}_t$  are to make first optimization steps more stable. Because the calculation of  $m_t$ ,  $v_t$  can be seen as geometric series



# Reduce On Plateau

Reduce learning rate by some factor if loss is not decreasing enough.


- we can't converge to exact local minima
- unfortunately, we increase "sensitivity" to more narrow local minima.





# Cycle LR

If weight initialization wasn't good enough, we try to increase learning rate at first few steps in a hope to jump into a better local minimum.



# Regularization

Most popular:

- 1  $L_2$  norm regularization through weight decay
- 2 Early stopping
- 3 Data augmentation. Create new samples from the same domain to increase size of your dataset. Remember generalization bounds.
- 4 Dropout. Drop random nodes in a layer with probability  $p$
- 5 Batch Normalization

# Dropout

There are 2 interpretations for dropout:

- "Bagging" over neural networks
- Avoid feature coadaptation

Difference between bagging and dropout:

$$p(y|x) = \frac{1}{K} \sum_{i=1}^K p_i(y|x)$$

for bagging

$$p(y|x) = \sum_{\mu} p(\mu) p(y|x, \mu)$$

for dropout, where  $\mu$  is mask on weights.

There is an exponential number of masks for fixed number of weights, that makes dropout more effective than explicit bagging.


# Dropout

**On training:** On each batch randomly remove neurons in the previous layer with probability  $p$ .


**On inference:**

Ideally, sample all  $2^n$  dropped-out networks and average predictions. In practice, approximate by using the full network with each node's output weighted by a factor of  $1 - p$ , so the expected value of the output of any node is the same as in the training stages.

=> Although it effectively generates  $2^n$  neural nets, but at test time only a single network needs to be tested.



(a) Standard Neural Net



(b) After applying dropout.

# BatchNorm

**On training:**  
on every batch  $t$ :

$$\mu_t = \frac{1}{m} \sum_{i=1}^m x_i$$

$$\sigma_t^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_t)^2$$

$$\hat{x}_i = \frac{x_i - \mu_t}{\sigma_t + \epsilon}$$

$$y_i = \gamma \hat{x}_i + \beta = \text{BN}_{\gamma, \beta}(x_i)$$

where  $\mu_t$  - estimated batch mean

$\sigma_t^2$  - estimated batch variance

$\hat{x}_i$  - normalized input

$\gamma$  - learnable scale parameter

$\beta$  - learnable shift parameter

# BatchNorm

**On inference:** we can't compute  $\mu_t, \sigma_t^2$ . Instead, we use some running average over  $\mu_t, \sigma_t^2$  that were observed during training.

- NN in theory can learn  $\gamma, \beta$  to undo batch normalization. In practice, they usually don't
- BatchNorm stabilizes training by making surface of loss function more smooth

